



Gene - Disease Association Prediction

Sistemas Inteligentes para a Bioinformática
Professor Miguel Rocha

Carlos Gomes (pg51681)
Diogo Esteves (pg2893)
Joana Lopes (pg53498)
Tiago Miranda (pg54437)

Overview



PROBLEM



DATA EXPLORATION AND
PRE-PROCESSING



MACHINE LEARNING
TASK



DEEP LEARNING TASK


Gene-Disease Association (GDA)

The study of **gene-disease associations** is a crucial step in understanding the etiology of diseases. Identifying the link between genes and diseases makes it possible to understand the disease's cause better and develop strategies to combat it. [1]

GDA quantify the relation among a pair of gene-disease and is one of the core concepts of **DisGeNET** platform that are integrated on TDC.

DisGeNET is a discovery platform containing one of the largest publicly available collections of genes and variants associated to human diseases. [2]

Using **machine learning** accurately can contribute to **new discoveries** and bring numerous **therapeutic opportunities**.



Task Description

Regression. Given the disease description and the amino acid sequence of the gene, predict their association.

GDA dataset : 52,476 gene-disease pairs (7 399 unique genes, 7095 unique diseases)

	Gene_ID	Gene	Disease_ID	Disease	Y
0	1	MSMLVVFLLLWGVTWGPVTEAAIFYETQPSLWAESESLLKPLANVT...	C0019209	Hepatomegaly: Abnormal enlargement of the liver.	0.30
1	1	MSMLVVFLLLWGVTWGPVTEAAIFYETQPSLWAESESLLKPLANVT...	C0036341	Schizophrenia: Schizophrenia is highly heritab...	0.30
2	2	MGKNKLLHPSLVLLLLVLLPTDASVSGKPQYMVLVPSLLHTETTEK...	C0002395	Alzheimer's Disease: Alzheimer disease is the ...	0.50
3	2	MGKNKLLHPSLVLLLLVLLPTDASVSGKPQYMVLVPSLLHTETTEK...	C0007102	Malignant tumor of colon: A primary or metasta...	0.31
4	2	MGKNKLLHPSLVLLLLVLLPTDASVSGKPQYMVLVPSLLHTETTEK...	C0009375	Colonic Neoplasms: A benign or malignant neopl...	0.30

Unique identifiers for genes.

Amino acids sequences corresponding to each geneID through the Uniprot platform.

Unique identifiers for diseases.

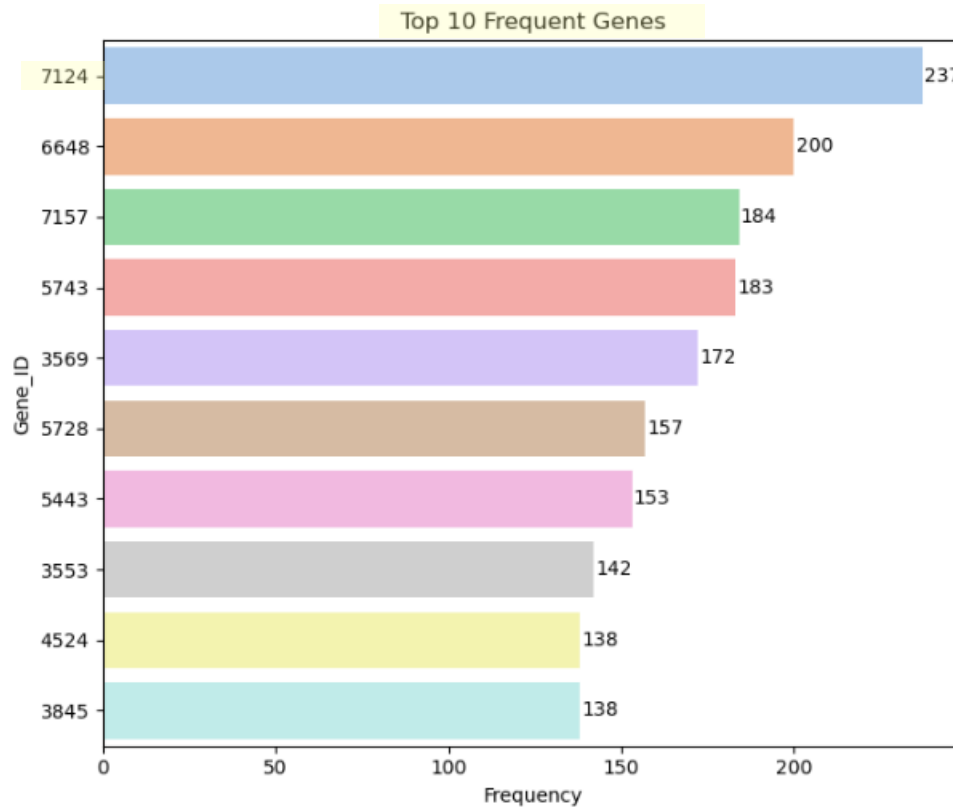
Disease definition corresponding to each diseaseID through the MedGen-NCBI platform.

Target: association score for a gene-disease pair.

```
data.print_stats()
```

```
--- Dataset Statistics ---  
7399 unique genes.  
7095 unique diseases.  
52476 gene-disease pairs.  
-----
```

Exploratory analysis: gene and disease frequency



These results indicate that there are **genes with greater involvement in multiple pathological conditions**. The high frequency of these genes may be a result of their **participation in fundamental biological processes**, such as cell cycle regulation, apoptosis or immune response.



GeneID	Description
7124	tumor necrosis factor

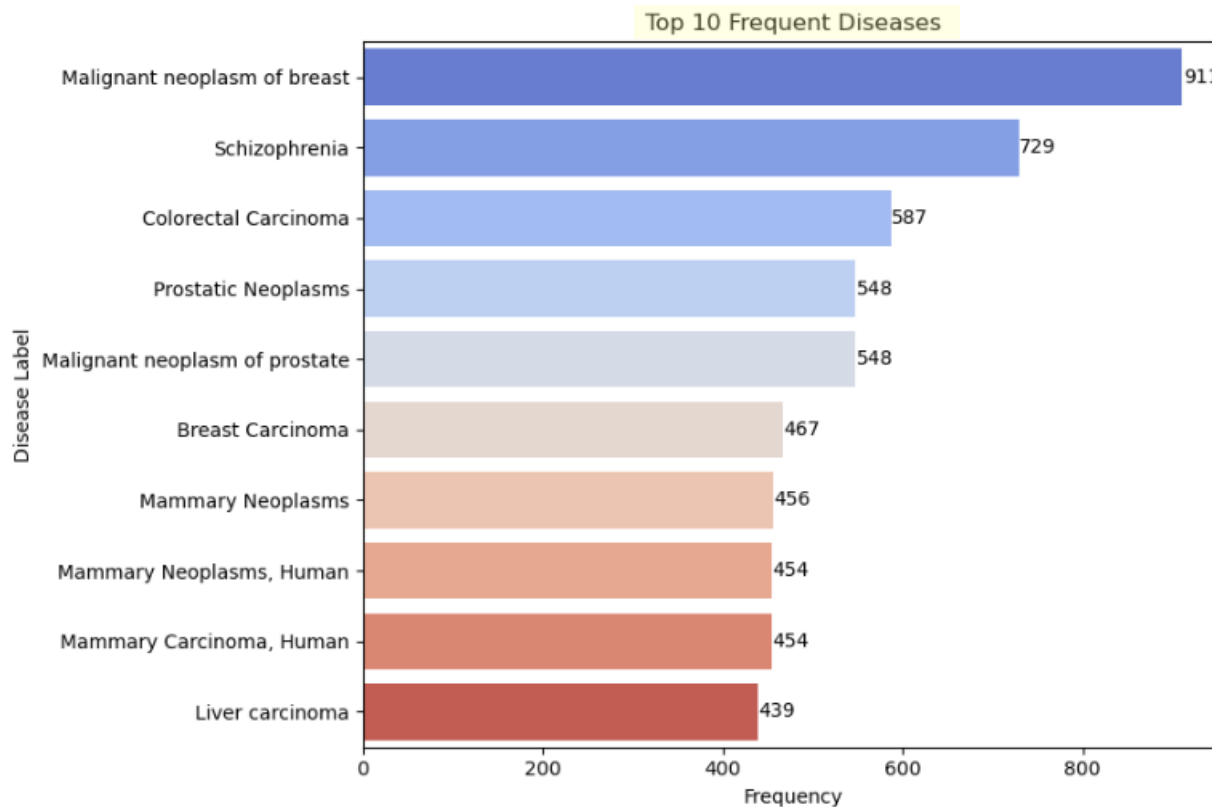
This gene encodes a **multifunctional proinflammatory cytokine** that belongs to the **Tumor Necrosis Factor** superfamily.

This cytokine is involved in the **regulation of a wide spectrum of biological processes** and has been **implicated in a variety of diseases**. [3]

Diseases associated with the TNF gene:

	Disease_Label
0	Adenocarcinoma
1	Albuminuria
2	Alzheimer's Disease
3	Anemia
4	Refractory anemias
5	Anemia, Sickle Cell
6	Anorexia
7	Anthraxis
8	Anxiety Disorders
9	Arthritis, Infectious

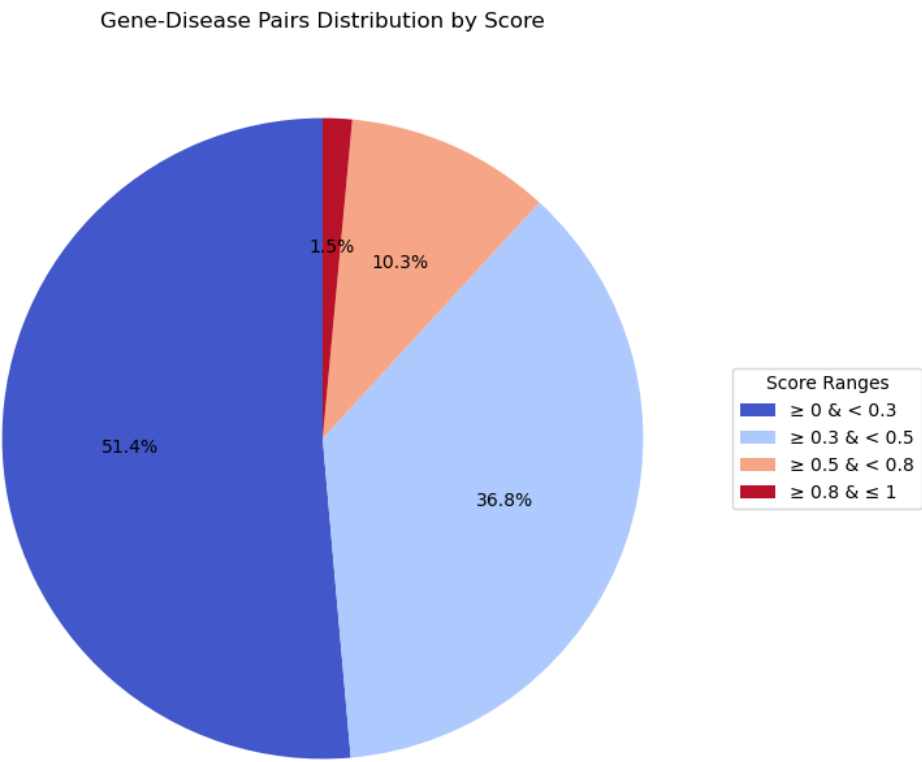
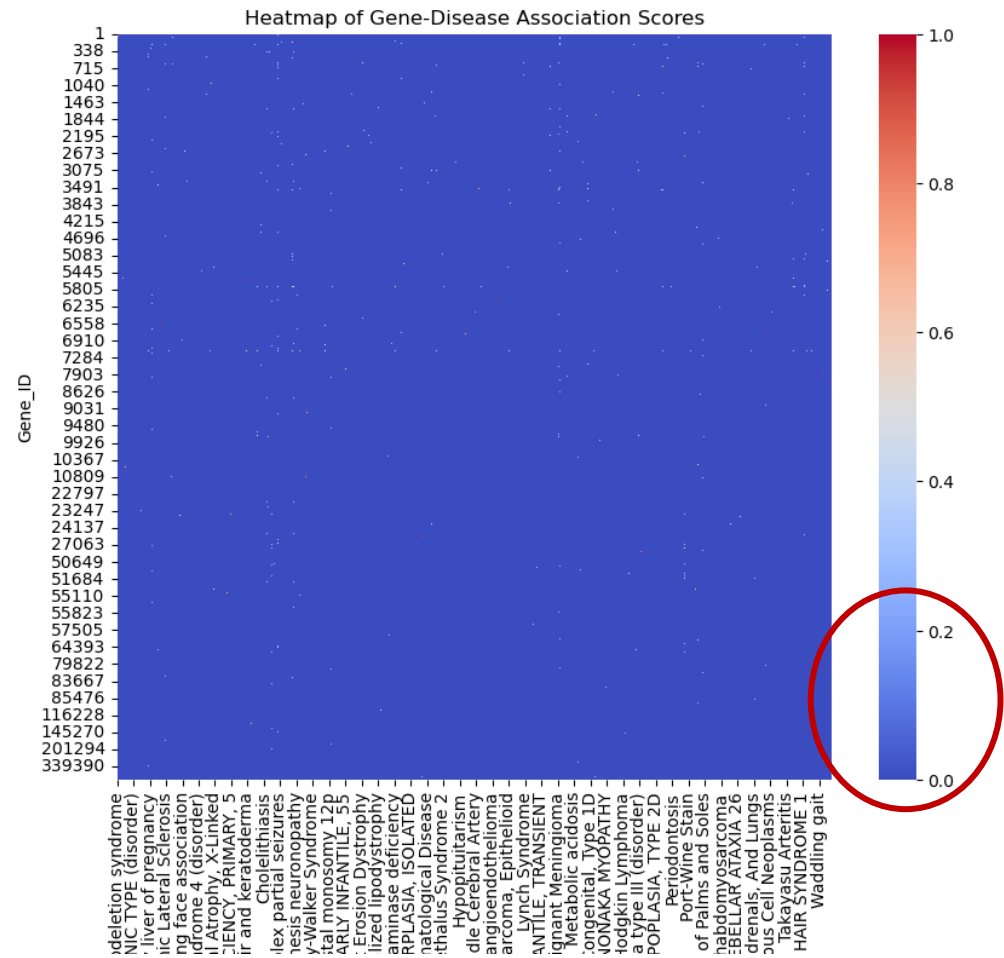
Exploratory analysis: gene and disease frequency



There is a clear pattern of highly frequent diseases in the dataset, led by conditions such as **malignant neoplasm of the breast** with 911 associations and **schizophrenia** with 729 associations.

The list is dominated by **neoplasms** such as breast, prostate and liver cancer, as well as the condition schizophrenia.

Exploratory analysis: pairs gene-disease score



Exploratory analysis: pairs gene-disease score

Number of Gene-Disease pairs with association score of 1: 335

Gene-Disease pairs with score = 1:

	Gene_ID	Disease_Label	Y
73	19	Tangier Disease	1.0
104	24	Stargardt's disease	1.0
106	24	STARGARDT DISEASE 1 (disorder)	1.0
163	34	Medium-chain acyl-coenzyme A dehydrogenase def...	1.0
169	35	Deficiency of butyryl-CoA dehydrogenase	1.0
178	37	Very long chain acyl-CoA dehydrogenase deficiency	1.0
400	90	Fibrodysplasia Ossificans Progressiva	1.0
887	175	Aspartylglucosaminuria	1.0
915	178	Glycogen Storage Disease Type III	1.0
1076	190	X-linked Adrenal Hypoplasia	1.0

Highly associated pairs can **guide drug development**, help **identify biomarkers** for **early diagnosis** and contribute to **personalizing treatments**.

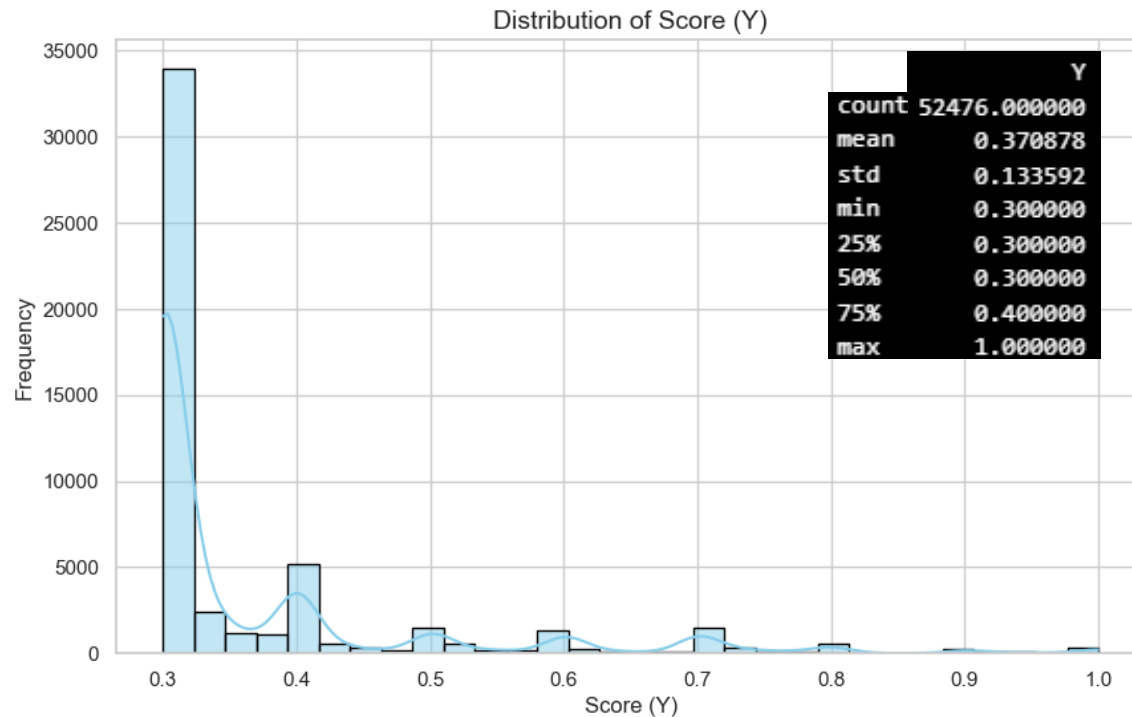


GeneID	Description
19	ATP binding cassette subfamily A member 1

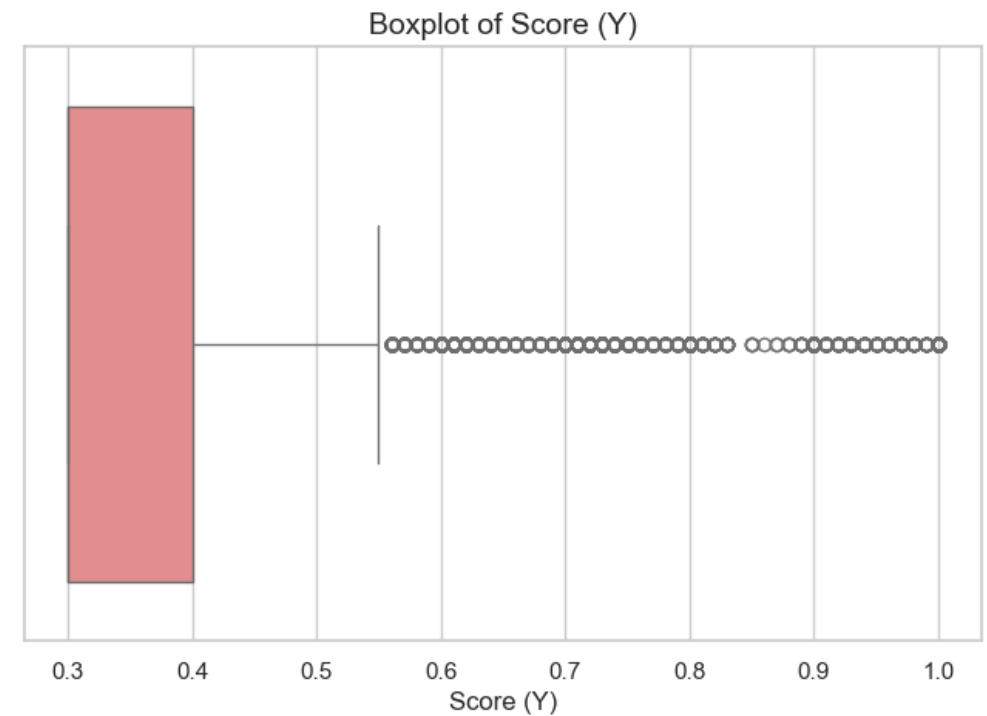
The membrane-associated protein encoded by this gene is a member of the **superfamily of ATP-binding cassette (ABC) transporters**. [4]

With cholesterol as its substrate, this protein functions as a **cholesterol efflux pump in the cellular lipid removal pathway**. Mutations in both alleles of this gene cause **Tangier disease** and **familial high-density lipoprotein (HDL) deficiency**. [5]

Exploratory analysis: Y distribution



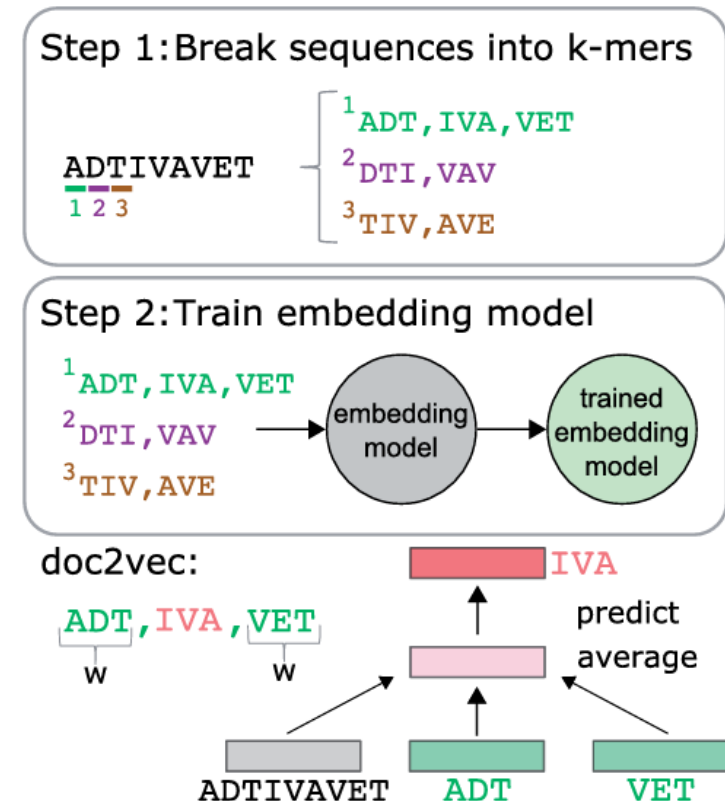
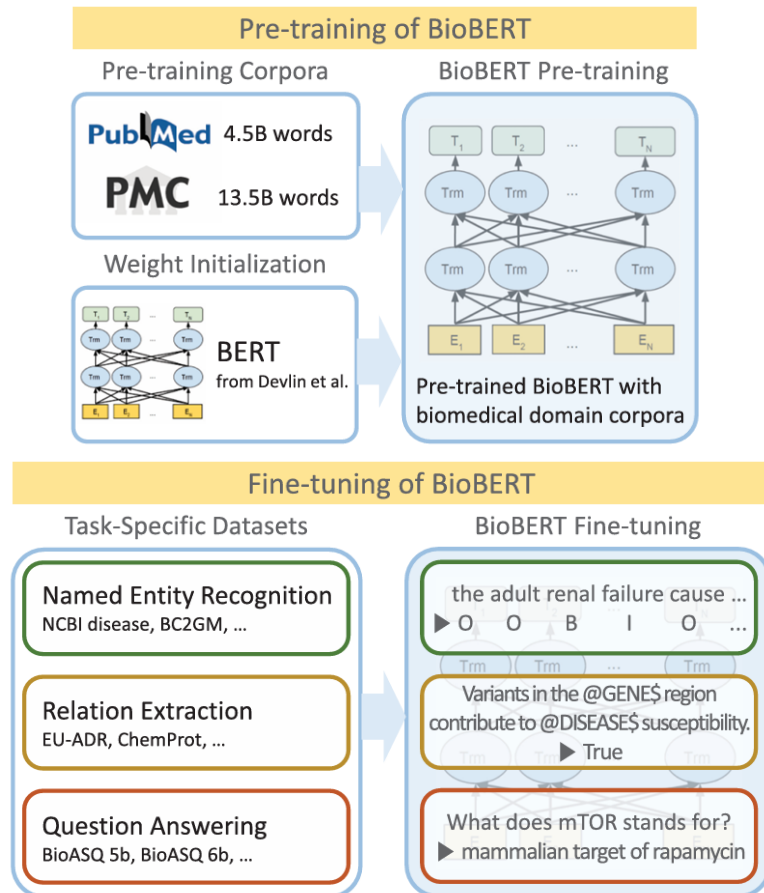
The target variable does not follow a normal distribution, with a clear **concentration of pairs in scores between 0.3 and 0.4**, indicate a low strength of association.



Based on the analysis of outliers using the interquartile range, **5475 outliers** were identified. In the context of our problem, these outliers represent gene-disease pairs where **a given gene contributes significantly to a given disease**.

Pre-processing task

- ✓ Transformation of the "Disease" Column (Textual Descriptions of Diseases): generate embeddings through **BioBERT** model [6]
- ✓ Transformation of the "Gene" Column (Amino Acid Sequences): generate embeddings through **k-mers + Word2Vec** [7]



Pre-processing task

- ✓ Conversion of the **Disease_Embedding** and **Seq_Embedding** columns from strings to **numerical arrays**.
- ✓ Feature integration: Combining generated features into a **unified dataset**.
- ✓ Data cleaning and normalization: **Handling missing values** and **scaling numerical features**.
- ✓ **Dataset splitting** dividing the data into training (60%), validation (20%), and test (20%) sets.

	Disease_Embedding	Seq_Embedding	Y
0	[0.90172192 0.90109823 0.88846408 0.91735898 0...	[0.58223012 0.98432244 0.92557997 0.5781199 0...	0.000000
1	[0.90579451 0.92227764 0.94480602 0.92059937 0...	[0.58223012 0.98432244 0.92557997 0.5781199 0...	0.000000
2	[0.90225344 0.88881582 0.89399212 0.91228476 0...	[0.58823403 1. 0.91926575 0.57808982 0...	0.285714
3	[0.88330894 0.91186235 0.8578483 0.8877738 0...	[0.58823403 1. 0.91926575 0.57808982 0...	0.014286
4	[0.89153572 0.89848732 0.90409447 0.91663919 0...	[0.58823403 1. 0.91926575 0.57808982 0...	0.000000
...
52471	[0.91774542 0.88312941 0.92196623 0.9114768 0...	[0.51241172 0.95092051 0.86576521 0.60375322 0...	0.014286
52472	[0.8771878 0.8883523 0.88011821 0.88788932 0...	[0.51241172 0.95092051 0.86576521 0.60375322 0...	0.142857
52473	[0.89576661 0.88151587 0.9057607 0.93911187 0...	[0.57023404 0.91414409 0.95015561 0.49892172 0...	0.000000
52474	[0.90726671 0.89780856 0.93095453 0.95476723 0...	[0.57023404 0.91414409 0.95015561 0.49892172 0...	0.000000
52475	[0.91257948 0.90366618 0.94029109 0.95948356 0...	[0.57023404 0.91414409 0.95015561 0.49892172 0...	0.000000

52476 rows x 3 columns

Train shape: (31485, 2), Validation shape: (10495, 2), Test shape: (10496, 2)

Dimension of Seq_Embedding embeddings:
Train: (100,)

Dimension of Disease_Embedding embeddings:
Train: (768,)

Concatenating features

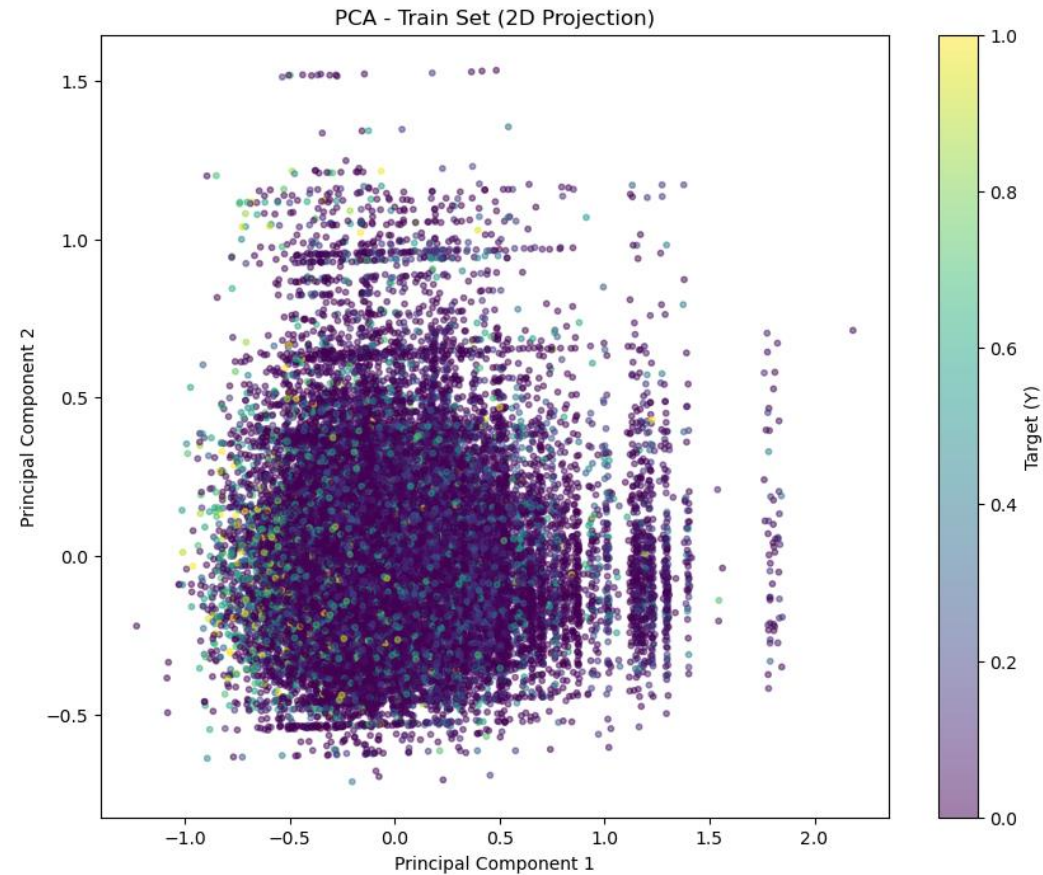
Shape of the feature matrix before PCA: (31485, 868)




MACHINE LEARNING

Unsupervised Learning

- After applying **PCA**, the original embeddings with 868 dimensions were reduced to **181 principal components**, preserving **95% of the explained variance**.



- **T-SNE** proved to be a **better choice for visualizing** the data.



Machine Learning task: main results

	Linear Regression		Ridge Regression			DecisionTreeRegressor			RandomForestRegressor			Suport Vector Regressor		
Hyperparameters	Def.	Def.	Def.	Def.	Randomized SearchCV	Def.	Def.	Randomized SearchCV	Def.	Def.	Randomized SearchCV	Def.	Def.	Randomized SearchCV
Reduced data by PCA	Yes	No	Yes	No	No	Yes	No	No	Yes	No	No	Yes	No	No
MSE	0,028	0,025	0,028	0,025	0,025	0,056	0,044	0,281	0,026	0,023	0,023	0,026	0,028	0,025
R-squared	0,229	0,306	0,229	0,306	0,311	-0,561	- 0,239	0,214	0,274	0,353	0,356	0,284	0,213	0,294

- ✓ Better performance of the models when using the **original data** (without applying PCA);
- ✓ Better performance of the models when using **RandomizedSearchCV** for **hiperparameter optimization**.

Best hyperparameters founded

n_estimators: 150
 min_samples_split: 5
 min_samples_leaf: 2
 max_features: sqrt
 max_depth: 30
 bootstrap: True

Most promising model

Machine Learning task: main results

	RFR
Hiperparameters	Randomized SearchCV
Reduced data by PCA	No
MSE	0,023
R-squared	0,357

```

Predictions with corresponding indices:
  Index Prediction
0      0  0.044039
1      1  0.004283
2      2  0.039293
3      3  0.121634
4      4  0.001997

Sample corresponding to the first prediction:
0      0.600617
1      1.000000
2      0.937280
3      0.593431
4      0.733264
...
863    0.908930
864    0.928018
865    0.894949
866    0.875315
867    0.897640
Name: 0, Length: 868, dtype: float64
    
```

```

Seq_Embedding:
0      0.600617
1      1.000000
2      0.937280
3      0.593431
4      0.733264
...
95     0.900016
96     0.550640
97     0.540475
98     0.492537
99     0.886927
Name: 0, Length: 100, dtype: float64

Disease_Embedding:
100    0.887657
101    0.895884
102    0.901867
103    0.911851
104    0.943052
...
863    0.908930
864    0.928018
865    0.894949
866    0.875315
867    0.897640
Name: 0, Length: 768, dtype: float64
    
```

	Gene_ID	Disease \
42741	23237	Alcoholic Intoxication, Chronic
42742	23237	Alzheimer's Disease
42743	23237	Presenile dementia
42744	23237	Movement Disorders
42745	23237	Schizophrenia
42746	23237	Cocaine-Related Disorders
42747	23237	Cocaine Dependence


```

Disease_Embedding
42741 [0.88765714 0.89588396 0.90186702 0.91185077 0...
42742 [0.90225344 0.88881582 0.89399212 0.91228476 0...
42743 [0.88421375 0.88678713 0.90185071 0.91114427 0...
42744 [0.88927211 0.92486154 0.8865273 0.88912351 0...
42745 [0.90579451 0.92227764 0.94480602 0.92059937 0...
42746 [0.90049218 0.92263105 0.92483915 0.91025301 0...
42747 [0.90432645 0.90552692 0.91940348 0.91358728 0...

Seq_Embedding
42741 [0.60061694, 1.0, 0.93728036, 0.59343133, 0.73...
42742 [0.60061694, 1.0, 0.93728036, 0.59343133, 0.73...
42743 [0.60061694, 1.0, 0.93728036, 0.59343133, 0.73...
42744 [0.60061694, 1.0, 0.93728036, 0.59343133, 0.73...
42745 [0.60061694, 1.0, 0.93728036, 0.59343133, 0.73...
42746 [0.60061694, 1.0, 0.93728036, 0.59343133, 0.73...
42747 [0.60061694, 1.0, 0.93728036, 0.59343133, 0.73...
    
```




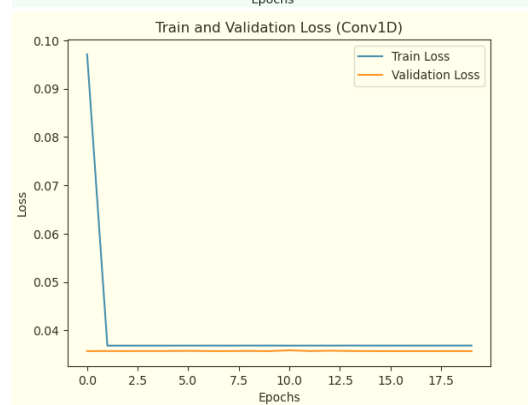

✓ Gene_ID: 23237

Activity regulated cytoskeleton associated protein [8]

Potential role in **vulnerability to addiction** [9]

Targeting Acr may represent a new mechanism for **preventing relapses related to chronic alcoholism** [10]

Deep Learning task : model metrics



	Dense NN	LSTM (v3)	LSTM Improved	Autoencoder	Conv1D
Validation Loss (MSE)	0.02	0.035	0.04	0.05	0.03
Train Loss (MSE)	0.015	0.033	0.037	0.045	0.02
Overfitting (val_loss - train_loss)	0.005	0.002	0.003	0.005	0.01
Generalization	High	Moderate	Moderate	Low	High
Stability	Stable	Slight Oscillations	Unstable	Unstable	Stable
Validation Loss (MSE)	0.02	0.035	0.04	0.05	0.03

✓ Better performances of the trained models;

Deep Learning task : model evaluation

	Dense NN	Conv1D	LSTM (v3)	LSTM (improved)	Autoencoder
Main Characteristics	General-purpose feedforward network, suited for non-sequential data.	Specialized for sequential data, effective for local patterns in time series.	Sequential network, capable of capturing long-term dependencies.	Incorporates Batch Normalization for more stable learning.	Learns compressed representations of data for reconstruction tasks.
Architecture	3 Dense layers (256 → 128 → 64) with Batch Normalization and Dropout.	2 Conv1D layers (64, 128 filters), MaxPooling1D, Flatten layer.	2 LSTM layers (32 → 16 units), Dense output layer.	2 LSTM layers with BatchNorm, followed by Dense (8 units).	Encoder (Dense 128 → 64), Decoder (Dense 128 → original dim).
Regularization	Dropout (30%/20%)	Dropout (30%)	Dropout (20%/10%)	Dropout (20%/10%)	BatchNorm, Sigmoid output
Validation Loss	0.02	0.03	0.035	0.04	0.05
Best Features	Effective for non-sequential data like embeddings	Ideal for sequential features in embeddings	Best of the 4 models for capturing gene-disease temporal relationships	Improves temporal dependencies with added stability	Best for dimensionality reduction in large datasets
Notes	Performs well with tabular data from gene-disease associations, good generalization	Stable loss; captures local patterns in sequential embeddings	Handles long-term dependencies in sequential embeddings	Overfitting observed; Used BatchNorm to help stabilize learning	Good for compressing embeddings but struggles with generalization

✓ Better performance – Dense NN, Conv1D

✓ High Potential – LSTM, improved LTSM



Future Work

- Final notebook adjustments.
- Optimize DL models, exploring a multimodal approach.

Thank you for your attention.

References

- [1] Opap, K., & Mulder, N. (2017). Recent advances in predicting gene–disease associations. *F1000Research*, 6, 578. <https://doi.org/10.12688/f1000research.10788.1>
- [2] Therapeutics Data Commons. (2024). TDC. <https://tdcommons.ai/>
- [3] TNF tumor necrosis factor [Homo sapiens (human)] - Gene - NCBI. (n.d.). *Www.ncbi.nlm.nih.gov*. <https://www.ncbi.nlm.nih.gov/gene/7124>
- [4] ABCA1 ATP binding cassette subfamily A member 1 [Homo sapiens (human)] - Gene - NCBI. (2020). *Nih.gov*. <https://www.ncbi.nlm.nih.gov/gene/19>
- [5] Koseki, M., Yamashita, S., Ogura, M., Ishigaki, Y., Ono, K., Tsukamoto, K., Hori, M., Matsuki, K., Yokoyama, S., & Harada-Shiba, M. (2021). Current Diagnosis and Management of Tangier Disease. *Journal of Atherosclerosis and Thrombosis*, 28(8), 802–810. <https://doi.org/10.5551/jat.rv17053>
- [6] Lee, J., Yoon, W., Kim, S., Kim, D., Kim, S., So, C. H., & Kang, J. (2019). BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4). <https://doi.org/10.1093/bioinformatics/btz682>
- [7] Yang, K. K., Wu, Z., Bedbrook, C. N., & Arnold, F. H. (2018). Learned protein embeddings for machine learning. *Bioinformatics*, 34(15), 2642–2648. <https://doi.org/10.1093/bioinformatics/bty178>
- [8] *ARC activity regulated cytoskeleton associated protein [Homo sapiens (human)]* - Gene - NCBI. (2025). *Nih.gov*. <https://www.ncbi.nlm.nih.gov/gene/23237>
- [9] Penrod, R. D., Thomsen, M., Taniguchi, M., Guo, Y., Cowan, C. W., & Smith, L. N. (2020). The activity-regulated cytoskeleton-associated protein, Arc/Arg3.1, influences mouse cocaine self-administration. *Pharmacology, Biochemistry, and Behavior*, 188, 172818. <https://doi.org/10.1016/j.pbb.2019.172818>
- [10] Pagano, R., Salamian, A., Zielinski, J., Beroun, A., Nalberczak-Skóra, M., Skonieczna, E., Cały, A., Tay, N., Banaschewski, T., Desrivères, S., Grigis, A., Garavan, H., Heinz, A., Brühl, R., Martinot, J.-L., Martinot, M.-L. P., Artiges, E., Nees, F., Orfanos, D. P., & Poustka, L. (2022). Arc controls alcohol cue relapse by a central amygdala mechanism. *Molecular Psychiatry*. <https://doi.org/10.1038/s41380-022-01849-4>